

Berner Fachhochschule - Technik und Informatik

Algorithms and Data Structures

Priority Queues and Heaps

Dr. Rolf Haenni

Fall 2008

Outline

Priority Queues

Heaps

Heap-Based Priority Queues

Bottom-Up Heap Construction

Outline

Priority Queues

Heaps

Heap-Based Priority Queues

Bottom-Up Heap Construction

Priority Queue ADT

- ▶ A priority queue stores a collection of items
- ▶ An **item** is a pair (*key*, *element*)
- ▶ Characteristic operations:
 - `insertItem(k,e)`: inserts an item with key *k* and element *e*
 - `removeMin()`: removes the item with the smallest key and returns its element
 - `minKey()`: returns the smallest key of an item (no removal)
 - `minElement()`: returns the element of an item with smallest key (no removal)
- ▶ General operations:
 - `size()`: returns the number of items
 - `isEmpty()`: indicates whether the priority queue is empty
- ▶ Two distinct items in a priority queue can have the same key

Total Order

- ▶ Keys in a priority queue can be arbitrary objects on which a **total order** is defined
- ▶ Mathematically, a total order is a **binary relation** \preceq defined on a set K which satisfies three properties:
 - **Totality**: $x \preceq y$ or $y \preceq x$, for all $x, y \in K$
 - **Antisymmetry**: $x \preceq y$ and $y \preceq x$ implies $x = y$, for all $x, y \in K$
 - **Transitivity**: $x \preceq y$ and $y \preceq z$ implies $x \preceq z$, for all $x, y, z \in K$
- ▶ Totality implies **Reflexivity**: $x \preceq x$, for all $x \in K$
- ▶ Examples: \leq or \geq for \mathbb{R} , alphabetical or reverse alphabetical order for $\{A, \dots, Z\}$, lexicographical order for $\{A, \dots, Z\}^*$, etc.

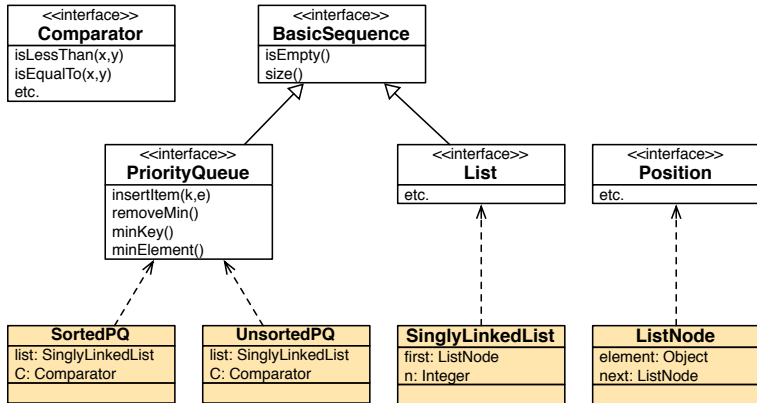
Comparator ADT

- ▶ A comparator encapsulates the action of comparing two keys according to a given total order relation
- ▶ A generic priority queue uses an auxiliary comparator (passed as a parameter to the constructor)
- ▶ When the priority queue needs to compare two keys, it uses its comparator
- ▶ Operations (all with Boolean return type):
 - `isLessThan(x,y)`
 - `isLessThanOrEqualTo(x,y)`
 - `isEqualTo(x,y)`
 - `isGreaterThan(x,y)`
 - `isGreaterThanOrEqualTo(x,y)`
 - `isComparable(x)`

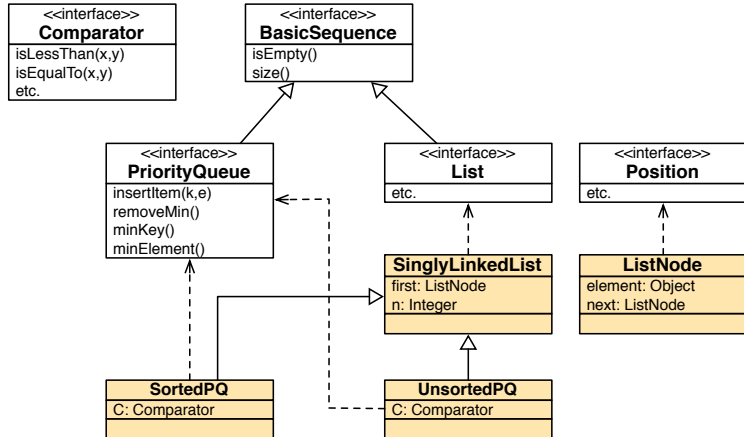
Sequence-Based Priority Queue

- ▶ There are two ways to implement a priority queue with a sequence (list or vector)
- ▶ Using an unsorted sequence
 - `insertItem(k,e)` runs in $O(1)$ time, since we can insert the item at the beginning of the sequence
 - `removeMin()`, `minKey()`, `minElement()` run in $O(n)$ time since we have to traverse the entire sequence to find the smallest key
- ▶ Using a sorted sequence
 - `insertItem(k,e)` runs in $O(n)$ time, since we have to find the place where to insert the item
 - `removeMin()`, `minKey()`, `minElement()` run in $O(1)$ time since the smallest key is at the beginning or end of the sequence

UML Diagram: Composition



UML Diagram: Inheritance



Sorting with a Priority Queue

- ▶ We can use a priority queue to sort a collection of comparable elements
 - Insert the elements one by one with a series of `insertItem(e,e)` operations
 - Remove the elements in sorted order with a series of `removeMin()` operations
- ▶ Depending on how the priority queue is implemented (sorted or unsorted), this leads to two well-known sorting algorithms
 - Selection-Sort
 - Insertion-Sort

PQ-Sort Algorithm

Solution in pseudo-code:

```
Algorithm PQ-Sort( $S, C$ )  
   $P \leftarrow$  new priority queue with comparator  $C$   
  while not  $S.isEmpty()$  do  
     $e \leftarrow S.removeElement(S.first())$   
     $P.insertItem(e, e)$   
  while not  $P.isEmpty()$  do  
     $e \leftarrow P.removeMin()$   
     $S.insertLast(e)$   
  return  $S$ 
```

Selection-Sort


- ▶ Selection-sort is the variation of PQ-sort where the priority queue is implemented with an **unsorted** sequence
- ▶ The main task is to **select** the elements from the unsorted sequence in the right order
- ▶ Running time analysis:
 - Inserting the elements into the priority queue with n `insertItem(e,e)` operations runs in $O(n)$ time
 - Removing the elements in sorted order from the priority queue with n `removeMin()` operations takes time proportional to $n + \dots + 2 + 1$, and thus runs in $O(n^2)$ time
- ▶ Selection-sort runs in $O(n^2)$ time

Insertion-Sort

- ▶ Selection-sort is the variation of PQ-sort where the priority queue is implemented with an **sorted** sequence
- ▶ The main task is to **insert** the elements at the right place
- ▶ Running time analysis:
 - Inserting the elements into the priority queue with n `insertItem(e,e)` operations takes time proportional to $1 + 2 + \dots + n$, and thus runs in $O(n^2)$ time
 - Removing the elements in sorted order from the priority queue with n `removeMin()` operations runs in $O(n)$ time
- ▶ In general (worst case), insertion-sort runs in $O(n^2)$ time
- ▶ If the input sequence is in reverse order (best case), it runs in $O(n)$ time

In-Place Sorting

- ▶ Instead of using external data structures, we can implement insertion- and selection-sort *in-place*
- ▶ A portion of the input sequence itself serves as priority queue
- ▶ For in-place insertion-sort we
 - keep the initial portion of the sequence sorted
 - use `swapElements(p,q)` instead of modifying the sequence



Outline

Priority Queues

Heaps


Heap-Based Priority Queues

Bottom-Up Heap Construction

The Heap ADT

- ▶ A heap is a specialized binary tree storing keys at its nodes (positions) and satisfying the following properties:
 - **Heap-Order**: $key(p) \succeq key(parent(p))$, for every node p other than the root
 - **Completeness**: let h be the height of the tree
 - 1) there are 2^i nodes of depth i for $i = 0, \dots, h - 1$
 - 2) at depth h , nodes are filled up from the left
- ▶ The last node of a heap is the rightmost node of depth h
- ▶ Heap operations:
 - `insertKey(k)`: inserts a key k into the heap
 - `removeMin()`: removes smallest key and returns it
 - `minKey()`: returns the smallest key (no removal)



Heap Example



- ▶ A heap storing n keys has height $O(\log n)$


Insertion to a Heap

- The insertion of a key k consists of 3 steps:
 - Add a new node q (the new last node)
 - Store k at q
 - Restore the heap order property (discussed next)





Upheap

- ▶ Algorithm upheap restores the heap-order property
 - Swap k along an upward path from the insertion node
 - Stop when k reaches the root or a node whose parent has a key smaller than or equal to k
 - Since the height of the heap is $O(\log n)$, upheap runs in $O(\log n)$ time
- ▶ `insertKey(k)` runs in $O(\log n)$ time




Removal from a Heap

- ▶ The removal algorithm consists of 4 steps:
 - Return the key of the root
 - Replace the root key with the key k of the last node
 - Delete the last node
 - Restore the heap order property (discussed next)




Downheap

- ▶ Algorithm downheap restores the heap-order property
 - Swap k along an downward path from the root (always with the smaller key of its children)
 - Stop when k reaches a leaf or a node whose children have keys greater than or equal to k
 - Since the height of the heap is $O(\log n)$, downheap runs in $O(\log n)$ time
- ▶ `removeMin()` runs in $O(\log n)$ time



Finding the Insertion Node


- ▶ Starting from the last node, the insertion node can be found by traversing a path of $O(\log n)$ nodes
 - If the last node is a left child, return the parent node
 - While the current node is a right child, go to the parent node
 - If the current node is a left child, go to the right child
 - While the current node is internal, go to the left child



Array-Based Heap Implementation

- ▶ We can represent a heap with n keys directly by means of an array of length $N > n$
 - By-pass the binary tree ADT
 - No explicit position ADT needed
- ▶ Idea similar to array-based implementation of binary trees
 - The array cell at index 0 is unused
 - The root of the heap is at index 1
 - The left child of the node at index i is at index $2i$
 - The right child of the node at index i is at index $2i + 1$
- ▶ The insertion operations corresponds to inserting at index $n + 1$ and thus runs in $O(1)$ time

Array-Based Heap: Example



Outline

Priority Queues

Heaps

Heap-Based Priority Queues

Bottom-Up Heap Construction

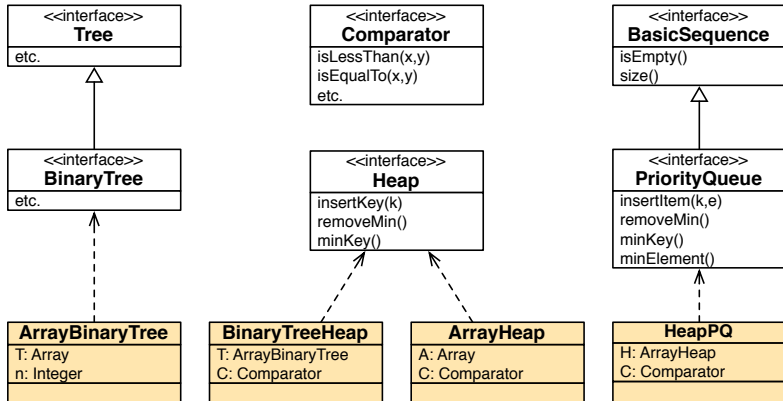
Heap-Based Priority Queues

- ▶ We can use a heap to implement a priority queue by storing a pair $item = (key, element)$ at each node of the heap
- ▶ Running times for different implementations

Operation	Unsorted Sequence	Sorted Sequence	Heap
size, isEmpty	1	1	1
minElement, minKey	n	1	1
insertItem	1	n	$\log n$
removeMin	n	1	$\log n$
PQ-Sort	n^2	n^2	$n \cdot \log n$

- ▶ In the long run, the heap-based implementation beats any sequence-based implementation

UML Diagram: Composition




Heap-Sort

- ▶ Heap-sort is the variation of PQ-sort where the priority queue is implemented with a heap
- ▶ Running time analysis:
 - Inserting the elements into the priority queue with n `insertItem(e,e)` operations runs in $O(n \log n)$ time
 - Removing the elements in sorted order from the priority queue with n `removeMin()` operations runs in $O(n \log n)$ time
- ▶ In general (worst case), heap-sort runs in $O(n \log n)$ time
- ▶ For large n , heap-sort is much faster than quadratic sorting algorithms such as insertion-sort or selection-sort

In-Place Heap-Sort

- ▶ To implement Heap-Sort *in-place*, use the front of the original array to implement the heap
 - Step 1: Use the reverse comparator (e.g. \succeq instead of \preceq) to build up the heap (by swapping elements)
 - Step 2: Iteratively remove the maximum element from the heap and insert it in front of the sequence



Outline

Priority Queues


Heaps

Heap-Based Priority Queues

Bottom-Up Heap Construction

Bottom-up Heap Construction

- ▶ We can construct a heap storing $n = 2^h - 1$ keys using a recursive bottom-up construction with $h - 1$ phases
- ▶ In phase i , pairs of heaps with $2^i - 1$ keys are merged into heaps with $2^{i+1} - 1$ keys
- ▶ Merging two heaps (and a new key k):
 - Create a new heap with the root node storing k and with the two heaps as subtrees
 - Perform downheap to restore the heap-order property




Recursive Algorithm in Pseudo-Code


```
Algorithm bottomUpHeap(S) // sequence of keys of length  $2^h - 1$ 
  if S.isEmpty() then
    return new heap
  else
     $k \leftarrow S.\text{elemAtRank}(0)$ 
    S.removeAtRank(0)
     $H_1 \leftarrow \text{bottomUpHeap}(\text{firstHalf}(S))$  // 1st recursive call
     $H_2 \leftarrow \text{bottomUpHeap}(\text{secondHalf}(S))$  // 2nd recursive call
    return mergeHeaps( $k, H_1, H_2$ ) // merge the results
```

In the general case, i.e. if the sequence is not of size $n = 2^h - 1$, we can adapt the splitting of the sequence accordingly


Example: Phase 1



Example: Phase 2



Example: Phase 3



Complexity Analysis

- ▶ The recursive `BottomUpHeap` algorithm decomposes a problem of size n into two problems of size $\frac{n-1}{2}$
 - Rule D2 with $q = r = 2$ (see Topic 2, Page 31)
- ▶ The running time $f(n)$ of each recursive step is $O(\log n)$
 - with an array-based sequence implementation, `firstHalf()` and `secondHalf()` run in $O(1)$ time
 - the necessary downheap in `mergeHeaps` runs in $O(\log n)$ time
- ▶ Apply the second column of D2, i.e. bottom-up heap construction runs in $O(n)$
- ▶ Speeds up the heap construction (Phase 1) of the Heap-Sort algorithm, but not Phase 2